

SADASHIV RAJ BHARADWAJ

Bangalore, India | +91-8660328606 | sadashivrajbharadwaj@gmail.com | LinkedIn

SUMMARY

Technical owner of Cisco's enterprise GenAI agent platform — **39 production agents, 17K users, 4 business units**. 9 years spanning LLM serving and evaluation, multi-agent orchestration, RAG, and petabyte-scale data infrastructure, including AdTech audience and marketing-attribution systems at Uber scale.

EXPERIENCE

Cisco Systems

Senior Software Engineer, Tech Lead — GenAI Platform

Bangalore, India

Jul 2020 – Present

- **Own the architecture of Cisco's multi-tenant GenAI agent platform**, serving Marketing, Sales, CX, and Splunk at **99.98% uptime**. Set technical direction across **4+ engineering teams** and hold final review on the platform-core path for all agent contributions.
- Designed **DeepAgent**, an **in-house** skills-based agentic runtime now the **org-wide standard adopted by 7+ teams**. Tool auto-discovery, sub-agent delegation, parallel LLM execution, and structured output parsing collapsed new-agent onboarding from **weeks to hours**, turning one-off bot projects into a self-service platform.
- Cut **LLM token spend ~50% platform-wide** with a two-tier inference strategy: prompt caching with 1-hour TTL plus context compaction that offloads long-session history onto a cheaper model class — holding cost flat as agent count and traffic grew.
- Standardized agent interoperability behind an **MCP and A2A layer**, exposing every agent as a callable service and enabling **cross-team multi-agent composition** in place of bespoke point-to-point integrations.
- Established **LLM evaluation and observability** as shared infrastructure — LangSmith trace-level debugging, automated eval harnesses (correctness, faithfulness, latency), and OpenTelemetry tracing — now covering **15 production agents** with regression detection their teams did not have to build.

Visibility AI — Ambient Executive Agent

- Conceived and launched an **ambient agent that saves Director-and-above leaders 6–8 hours per week**, replacing manual status-gathering across their project portfolios with one mobile Webex surface federating **6 MCP servers** — Asana, Microsoft 365 (Outlook/SharePoint), Airtable, and Webex messaging/meetings — under **per-user OAuth**.
- Implemented the **proactive digest engine**: webhook- and poll-driven change detection with watermark de-duplication, strict per-user token isolation (forged payloads are no-ops by design), and cost-tiered escalation from templated cards to LLM summaries — running **daily for 13 executives with zero recurring user interaction**.
- Shipped **brand-governed document generation** (PPTX/XLSX/DOCX/PDF) driven by a **Deck Designer sub-agent** that plans deck structure in an isolated context window, then composes slides under machine-enforced brand constraints — turning hours of manual assembly into a single prompt.

Marketing AI — Text-to-SQL over Marketing Performance Data

- Built a Text-to-SQL agent over **77 Snowflake tables / 2,255 columns** spanning campaigns, paid media, fractional attribution, and pipeline — with automated schema-to-skill generation, fiscal-calendar reasoning, group-based data authorization, and derived-metric and mandatory-filter validation that **eliminated silent wrong answers**. Reduced ad-hoc analyst SQL workload **~40%**.
- Added a **hybrid RAG retrieval layer** (vector + keyword over Weaviate) with chunking, re-ranking, and context-window optimization, now reused across other agents' knowledge bases.

Earlier at Cisco — Multimodal AI & Data Platform

- Delivered an end-to-end **multimodal pipeline** (ASR + speaker diarization + LLM summarization) turning executive conference video into transcripts, structured summaries, and embeddings — **4–8 hours to under 10 minutes**.
- Co-architected a YAML-driven distributed ETL framework that **replaced a \$100K/yr licensed tool** and cut data-product delivery from 7 days to 4 hours; reduced AWS spend **\$250K/yr (60% runtime reduction)** via multi-threaded shared-pool schedulers.

Gartner

Data Engineer

Gurgaon, India

Jan 2020 – Jul 2020

- Migrated legacy ingestion to Apache Airflow on GCP, **saving \$65K/yr**, and stood up 40+ high-frequency integrations delivering near-real-time session data into BigQuery.

Nineleaps Technology (contracted to Uber)

Bangalore, India

Software Development Engineer II

Sep 2017 – Jan 2020

- Re-engineered **Uber's AdTech audience pipeline** in PySpark, cutting processing **94% (6 hrs → 20 min)** and unlocking same-day campaign launches across markets.
- Optimized marketing spend allocation over large-scale rider/driver datasets in Presto and Hive, delivering a measured **10% lift in marketing ROI**.

TECHNICAL SKILLS

GenAI & LLM Systems: Agent runtime design, multi-agent orchestration, RAG, Text-to-SQL, prompt engineering & caching, context compaction, model evaluation, MCP, A2A | Claude (AWS Bedrock), OpenAI, Azure OpenAI, Gemini | Weaviate, LangGraph, LangSmith
Infrastructure: Kubernetes, Docker, FastAPI, Celery, Redis, gRPC, OpenTelemetry, OAuth 2.0, Jenkins CI/CD
Data & Cloud: Spark, Kafka, Airflow, Snowflake, BigQuery, Presto/Hive, PostgreSQL | GCP (BigQuery, Dataproc, GCS), AWS (Bedrock, EMR, Lambda, S3)
Languages: Python, SQL, Scala, Go, JavaScript, Bash

EDUCATION

University of Delhi — Bachelor of Technology, Electronics Engineering

Delhi, India | 2013 – 2017